

CLASS PROTOTYPE-BASED CLEANER FOR LABEL NOISE LEARNING

Anonymous authors

Paper under double-blind review

ABSTRACT

Semi-supervised learning based methods are current SOTA solutions to the noisy-label learning problem, which rely on learning an unsupervised label cleaner first to divide the training samples into a clean labeled set and a noisy unlabeled set. Typically, the cleaner is obtained via fitting a mixture model to the distribution of per-sample training losses. However, the modeling procedure is *class agnostic* and assumes the loss distributions of clean and noisy samples are the same across different classes. Unfortunately, in practice, such an assumption does not always hold due to the varying learning difficulty of different classes, thus leading to sub-optimal label noise partition criteria. In this work, we first reveal this long-ignored problem and propose a simple yet effective solution, named **Class Prototype-based label noise Cleaner (CPC)**. Unlike previous works treating all the classes equally, CPC fully considers loss distribution heterogeneity and applies class-aware modulation to partition the clean and noisy data. CPC takes advantage of loss distribution modeling and intra-class consistency regularization in feature space simultaneously and thus can better distinguish clean and noisy labels. We theoretically justify the effectiveness of our method by explaining it from the Expectation-Maximization (EM) framework. Extensive experiments are conducted on the noisy-label benchmarks CIFAR-10, CIFAR-100, Clothing1M and WebVision. The results show that CPC brings about impressive performance improvement across all benchmarks. Moreover, CPC shows outstanding performance especially in the extremely noisy scenarios, and improves the accuracy on CIFAR-100 at 90% noise rate by as high as 13% over the SOTAs.

1 INTRODUCTION

Deep Neural Networks (DNNs) have brought about significant progress to the computer vision community over past few years. One key to its success is the availability of large amount of training data with proper annotations. However, label noise is very common in real-world applications. Without proper intervention, DNNs would be easily misled by the label noise and yield poor performance.

In order to improve the performance of DNNs when learning with noisy labels, various methods have been developed (Liu et al., 2020; Li et al., 2020a; Reed et al., 2014; Nishi et al., 2021). Among them, semi-supervised learning based methods (Nishi et al., 2021; Li et al., 2020a) achieve the most competitive results. The semi-supervised learning methods follow a two-stage pipeline. They first model the loss distribution of training samples to construct a noise cleaner based on the “small-loss prior” (Han et al., 2020), which says in the early stage of training, samples with smaller cross-entropy losses are more likely to have clean labels. The prior is widely adopted and demonstrated to be highly effective in practice (Han et al., 2020). Given the noise cleaner, the training samples are divided into a labeled clean set and an unlabeled noisy set. Then, semi-supervised learning strategies like MixMatch (Berthelot et al., 2019) are employed to train DNNs on the divided dataset.

The key to their performance lies in the accuracy of the label-noise cleaner (Cordeiro et al., 2022). Usually, a single Gaussian Mixture Model (GMM) (Li et al., 2020a) is used to model the loss distribution of all the training samples across different categories. However, this modeling procedure is class-agnostic, which assumes a DNN model has the *same learning speed* to fit the training samples in different categories, thus the same loss value on samples in different categories can reflect the same degree of noise likelihood.

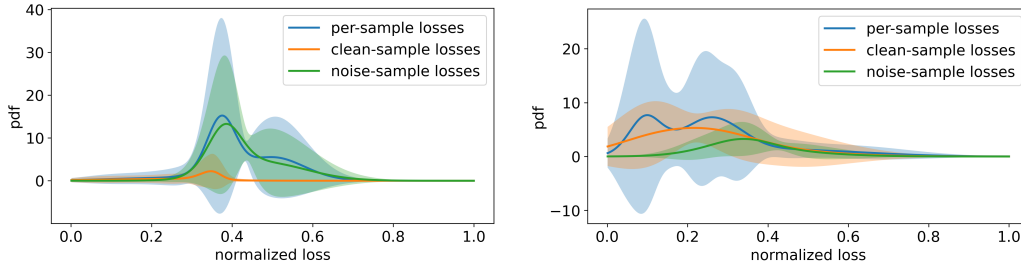


Figure 1: Loss distribution of samples in CIFAR-100 with 90% symmetric noise at epoch 30 (left) and CIFAR-10 with 40% asymmetric noise at epoch 10 (right), where the curves indicate mean probability density over all the categories while the shadow indicates the 95% confidence interval. The loss distribution for each class deviates significantly from the average loss distribution.

Unfortunately, such assumption does not hold in practise. In Fig. 1, we present the cross-entropy loss distribution of training samples in CIFAR-100 under 90% symmetric noise (left) and CIFAR-10 under 40% asymmetric noise (right). We conduct Kolmogorov-Smirnov test (Massey Jr, 1951) to quantify the loss distribution difference between the samples in each class and samples in the whole dataset. The results show that for 54% categories in CIFAR-100 under 90% symmetric noise, the loss distribution of their clean samples hold a p-value lower than 0.05¹, while the number for the distribution of noise samples is 53%. Therefore, the class-agnostic label noise cleaner, which establishes a overly rigid criterion shared by all the classes, would introduce more noise samples to the clean set while reject clean samples, and consequently get the model perform poorly. A straightforward remedy to the problem is to fit distinct GMMs to losses of samples in different classes respectively, yielding a class-aware GMM cleaner. Nevertheless, this class-aware modeling strategy implicitly assumes that label noise is existed in every class. In the case of asymmetric noise *e.g.*, CIFAR10-asym40%, where samples in parts of classes are clean, such a naive strategy would classify most of hard samples in the clean classes as noise, and results in negative affect on model training.

Considering that images in the same category should share similar visual representations, the distance between a sample and the cluster center (*e.g.*, class prototype) of its labeled class is helpful for recognizing label noise. In this paper, we propose a simple **Class Prototype-based label noise Cleaner (CPC)** to apply class-aware modulation to the partitioning of clean and noisy data, which takes advantage of intra-class consistency regularization in feature space and loss distribution modeling, simultaneously. CPC learns embedding for each class, *i.e.*, class prototypes, via intra-class consistency regularization, which urges samples in the same class to gather around the corresponding class prototype while pushes samples not belonging to the class away. Unlike the aforementioned naive class-aware GMM cleaner, CPC apply class-aware modulation to label noise partitioning via representation similarity measuring, which is more general for different label noise scenarios. Meanwhile, CPC leverages the “small-loss prior” to provide stronger and more robust supervision signals to facilitate the learning of prototypes.

We plug CPC to the popular DivideMix(Li et al., 2020a) framework, which iterates between label noise partitioning and DNNs optimization. With the stronger label noise cleaner in the first stage, DNNs can be trained better in the second stage, which would further improve the learning of class prototypes. We theoretically justify the procedure from Expectation-Maximization algorithm perspective, which guarantees the efficacy of the method.

We conduct extensive experiments on multiple noisy-label benchmarks, including CIFAR-10, CIFAR-100, Clothing1M and WebVision. The results clearly show that CPC effectively improves accuracy of label-noise partition, and brings about significant performance improvement across all noise levels and benchmarks. Moreover, CPC shows outstanding performance especially in an extremely noisy condition, which improves the test accuracy on CIFAR-100 with 90% noise by 13%.

The contribution of our work lie in three folds: (1) We reveal the long-ignored problem of *class-agnostic* loss distribution modeling that widely existed in label noise learning, and propose a simple

¹A p-value < 0.05 suggests that the class-wise loss distribution differ from the global loss distribution with a confidence higher than 95%.

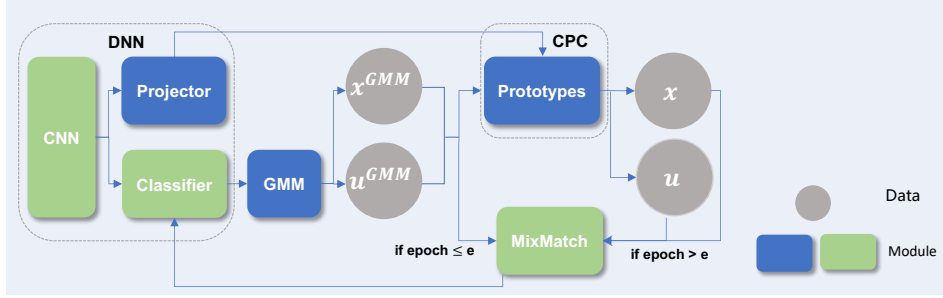


Figure 2: Illustration of the training pipeline in a single epoch. Blue modules are utilized in the first stage, where we update the prototypes in CPC and partition the training data. Green modules are utilized in the second stage, where the DNN model is trained based on the partitioned data.

yet effective solution, named Class Prototype-based label noise Cleaner (CPC); (2) **CPC takes advantage of loss distribution modeling and intra-class consistency regularization in feature space simultaneously**, which can better distinguish clean and noise labels; (3) Extensive experimental results show that our method achieves competitive performance compared to current SOTAs.

2 RELATED WORK

Recent advances in robust learning with noisy labels can be roughly divided into three groups. (a) **Label correction methods** aim to translate wrong labels into correct ones. Early studies rely on an auxiliary set with clean samples for clean label inference (Xiao et al., 2015a; Vahdat, 2017; Li et al., 2017b; Lee et al., 2018). Recent efforts focus on performing label correction procedures without supervision regarding clean or noisy labels. (Yi & Wu, 2019a; Tanaka et al., 2018) propose to jointly optimize labels during learning model parameters. Li et al. (2020b) propose to correct corrupted labels via learning class prototypes and utilize the pseudo-label generated by measuring the similarity between prototypes and samples to train model. Wu et al. (2021) and Li et al. (2021) introduce neighbouring information in feature space to correct noise label, and propose a graph-based method and a class prototype-based method, respectively. (b) **Sample selection methods** select potential clean samples for training to eliminate the effect of noisy labels on learning the true data distribution. (Han et al., 2018; Jiang et al., 2018; 2020; Yu et al., 2019) involve training two DNNs simultaneously and focus on the samples that are probably to be correctly labeled. (c) **Semi-supervised learning methods** conceal noisy labels and treat these samples as unlabeled data (Ding et al., 2018). DivideMix (Li et al., 2020a) is a typical algorithm among these works, which compromises an unsupervised label noise cleaner that divides the training data to a labeled clean set and an unlabeled noise set, followed by semi-supervised learning that minimize the empirical vicinal risk of the model. Inspired by DivideMix, a series of methods (Cordeiro et al., 2022; Nishi et al., 2021; Cordeiro et al., 2021) are proposed, which achieve SOTA performance. However, all these methods rely on the class-agnostic loss distribution modeling to achieve the label noise cleaner, which hinders the performance of the model. In our method, we propose the simple yet effective class prototype-based label noise cleaner to solve the problem. Besides, compared to previous prototype-based label noise learning methods (Li et al., 2020b; 2021), our method are different from them in two folds: (1) we utilize **prototypes as label noise cleaner** to effectively improve the semi-supervised learning based methods; (2) CPC takes advantage of both loss distribution modeling and intra-class consistency regularization in feature space simultaneously which learns better prototypes.

3 PRELIMINARY

In label noise learning, given a training set $D = (X, Y) = \{(x_i, y_i)\}_{i=1}^N$, where x_i is an image and $y_i \in \{1, 2, \dots, K\}$ is the annotated label over K classes, the label y_i could differ from the unknown true label \hat{y}_i . In this paper, we follow the popular label noise learning framework DivideMix (Li et al., 2020a), which first warms up the model for a few epochs by training on all the data using the standard cross-entropy loss, and then trains the model by iterating a two-stage pipeline. The pipeline comprises an unsupervised label cleaner Q to divide training samples into a clean labeled

set \mathcal{X} and a noisy unlabeled set \mathcal{U} , followed by a semi-supervised learning stage that trains the model to minimise the empirical vicinal risk (EVR) (Zhang et al., 2017):

$$\ell_{EVR} = \frac{1}{|\mathcal{X}'|} \sum_{(x'_i, y'_i) \in \mathcal{X}'} \ell_{\mathcal{X}'}(p(\tilde{y}'_i|x'_i), y'_i) + \frac{\lambda}{|\mathcal{U}'|} \sum_{(x'_i, y'_i) \in \mathcal{U}'} \ell_{\mathcal{U}'}(p(\tilde{y}'_i|x'_i), y'_i), \quad (1)$$

where \mathcal{X}' and \mathcal{U}' indicate MixMatch (Berthelot et al., 2019) augmented clean and noise set. $\ell_{\mathcal{X}'}$ and $\ell_{\mathcal{U}'}$ denote the losses for set \mathcal{X}' and \mathcal{U}' , which are weighted by λ . $p(\tilde{y}'_i|x'_i)$ is the softmax output of DNNs, where \tilde{y}'_i is the predicted label. For more details about EVR, please refer to the appendix A.1.

In Li et al. (2020a), the unsupervised label cleaner is operated under the “small-loss prior”, which is widely adopted and demonstrated to be highly effective (Han et al., 2020). The prior assumes that in the early stage of training, samples with smaller cross-entropy losses are more likely to have clean labels. The well known insight behind the “small-loss prior” is that DNNs tend to learn simple patterns first before fitting label noise (Arpit et al., 2017). Given a training sample x_i and the softmax output $p(\tilde{y}_i|x_i)$ of DNNs, where \tilde{y}_i is the predicted label, the cross-entropy loss $l(p(\tilde{y}_i|x_i), y_i)$ reflects how well the model fits the training sample.

To achieve the unsupervised label cleaner Q , a two-component Gaussian Mixture Model (GMM) is employed to fit the loss distribution of all training samples, i.e., $\ell(p(\tilde{y}_i|x_i), y_i) \sim \phi_0\mathcal{N}(\mu_0, \sigma_0) + \phi_1\mathcal{N}(\mu_1, \sigma_1)$, where $\mu_0 < \mu_1$, and ϕ is a mixing coefficient. The component with smaller mean represents the distribution of clean samples and the other one is for noise samples. We use $z_i \in \{0, 1\}$ indicates the data is clean or not. Then, $q(z_i = 0)$ represents the clean probability of x_i , which is the posterior probability of its loss belonging to the clean component. The label cleaner is shared by training samples across different classes, which is actually *class-agnostic*. A hypothesis implicitly accompanying this loss distribution modeling method is ignored by current works, which assumes the loss distributions of clean and noise samples are consistent across different categories. Unfortunately, as illustrated in Fig.1, the hypothesis dose not hold in practise. In this paper, we propose the class prototype-based label noise cleaner which applies class-aware modulation to the partitioning of clean and noisy data and improves label noise learning.

4 METHODOLOGY

In this section, we first introduce the overview of our method for learning with label noise. Then, we present the key component of our method—Class Prototype-based label noise Cleaner (CPC) in details. At last, we theoretically justify the efficacy of CPC from the Expectation Maximization algorithm perspective.

4.1 OVERVIEW

Our method follows the two-stage label noise learning framework DivideMix (Li et al., 2020a) and improves the framework with the proposed CPC. CPC comprises class prototypes $C = \{c_k \in \mathbb{R}^{1 \times d} | k = 1, 2, \dots, K\}$, where c_k indicates the prototype of k -th class and d is the dimension of prototype embedding. Our DNN model consists of a CNN backbone, a classifier head and a projection layer. The backbone maps an image input x_i to a feature vector $v_i \in \mathbb{R}^{1 \times D}$. The classifier takes v_i as input and outputs class prediction $p(\tilde{y}_i|x_i)$. The projection layer serves to project the high dimension feature v_i to a low-dimensional embedding $v'_i \in \mathbb{R}^{1 \times d}$, where $d < D$.

As shown in Fig. 2, we update the DNN as well as the CPC by iterating a two-stage training pipeline in every epoch. In the first stage, we update CPC and utilize the updated CPC to partition label noise. We first calculate the cross-entropy loss of every training sample and fits a GMM to the losses. We utilize the GMM as a label noise cleaner to get a labeled clean set \mathcal{X}^{GMM} and a unlabeled noise set \mathcal{U}^{GMM} . The data partition \mathcal{X}^{GMM} and \mathcal{U}^{GMM} are utilized to update the prototypes in CPC. Then, the updated CPC is employed to re-divide the training data into another two set \mathcal{X} and \mathcal{U} . In the second stage, we train DNN model to minimise the EVR in Eq. (11) with data partitioned by the cleaner. In the first e epochs, we wait CPC to warm up, and minimise the EVR of DNNs based on training data partitioned by the GMM cleaner. After the e -th epoch, the label noise estimation results of CPC, i.e., \mathcal{X} and \mathcal{U} are employed to train DNNs, while the estimation results of GMM cleaner are only used to update prototypes in CPC.

4.2 CLASS PROTOTYPE-BASED LABEL NOISE CLEANER

In order to apply class-aware modulation to the label noise partitioning, we propose to learn an embedding space where samples from the same class gather around its class prototypes, and leverage the prototypes to recognize noisy labels. The prototypes are typically learnt with intra-class consistency regularization, which urges samples in the same class to gather around the corresponding class prototype while pushes samples not belonging to the class away. Previous methods (Wang et al., 2022; Li et al., 2020b) apply the intra-class consistency regularization to prototype learning via unsupervised contrastive objectives, *e.g.*, prototypical contrastive objective (Li et al., 2020c). Compared to the unsupervised contrastive objectives whose performance highly depends on the quality of representation learnt by the CNN encoder, the GMM cleaner which is operated under the well evaluated “small-loss prior” can provide stronger and more robust supervision signals, especially in the early stage of training.

Therefore, we propose to take samples in clean set \mathcal{X}^{GMM} as positive samples and those in noise set \mathcal{U}^{GMM} as negative samples to update prototypes. Specifically, given the feature embedding v'_i of a sample $x_i \in \mathcal{X}^{\text{GMM}}$, we update prototypes C as well as the parameters of the projector to minimize the distance between $\langle c_{k=y_i}, v'_i \rangle$ and maximize the distance between $\langle c_{k \neq y_i}, v'_i \rangle$ via minimize the loss function $L_{\mathcal{X}^{\text{GMM}}}$:

$$L_{\mathcal{X}^{\text{GMM}}} = -\frac{1}{|\mathcal{X}^{\text{GMM}}|} \sum_{\mathcal{X}^{\text{GMM}}} \sum_{k=1}^K \ell_k(v'_i, y_i), \text{ where} \quad (2)$$

$$\ell_k(v'_i, y_i) = \begin{cases} \log(\text{sigmoid}(v'_i c_k^\top)), & k = y_i, \\ \lambda^{\text{neg}} \log(1 - \text{sigmoid}(v'_i c_k^\top)), & k \neq y_i, \end{cases}$$

where $\lambda^{\text{neg}} = \frac{1}{K}$ weights the losses between positive pair and negative pairs to avoid under-fitting the positive samples. Given v'_i of a sample $x_i \in \mathcal{U}^{\text{GMM}}$, we update prototypes c_k as well as the parameters of the projector to maximize the distance between $\langle c_{k=y_i}, v'_i \rangle$ via minimizing the loss function $L_{\mathcal{U}^{\text{GMM}}}$:

$$L_{\mathcal{U}^{\text{GMM}}} = -\frac{1}{|\mathcal{U}^{\text{GMM}}|} \sum_{\mathcal{U}^{\text{GMM}}} \log(1 - \text{sigmoid}(v'_i c_k^\top)), \text{ where } k = y_i. \quad (3)$$

At last, we collect training samples \mathcal{X}^P with predicted pseudo-label $\max_k(p(\tilde{y}_i|x_i))$ from \mathcal{U}^{GMM} , where $p(\tilde{y}_i|x_i)_k > \text{average}(\{p(\tilde{y}_i|x_i)_k | (x_j, y_j | y_j = k) \in \mathcal{X}^{\text{GMM}}\})$, and update the prototypes and projectors to minimize the loss function $L_{\mathcal{X}^P}$:

$$L_{\mathcal{X}^P} = -\frac{1}{|\mathcal{X}^P|} \sum_{\mathcal{X}^P} \log(\text{sigmoid}(v'_i c_k^\top)), \text{ where } k = \max_k(p(\tilde{y}_i|x_i)). \quad (4)$$

The overall loss function L_C for prototypes and the projector is as follows:

$$L_C = L_{\mathcal{X}^{\text{GMM}}} + L_{\mathcal{U}^{\text{GMM}}} + \alpha L_{\mathcal{X}^P}, \quad (5)$$

where α is the weight scalar.

With prototypes C , CPC distinguishes a clean sample (x_i, y_i) with the score $q(z_i = 0) = \text{sigmoid}(v'_i c_{k=y_i}^\top)$ and the threshold τ . Samples with $q(z_i = 0) > \tau$ are classified as clean, and otherwise as noise.

4.3 THEORETICAL JUSTIFICATION ON THE EFFICACY OF CPC

We provide theoretical justification on the efficacy of CPC from the perspective of Expectation-Maximization algorithm, which guarantees that though CPC does not follow the classical prototypical contrastive objective, it can still learn meaningful prototypes and act as an effective cleaner.

We consider training data with label noise $D = (X, Y) = (x_i, y_i)_{i=1}^N$ as the observable data, and $Z \in \{0, 1\}^N$ as the latent variable, where $z_i = 0$ iff (x_i, y_i) is clean (*i.e.*, $y_i = \hat{y}_i$). The prototypes C in the cleaner are taken as parameters expected to be updated. Then, the negative log likelihood

for D given C is as follows:

$$NLL(D|C) = - \sum_D \log \sum_{z_i \in \{0,1\}} p(x_i, y_i, z_i|C) = - \sum_D \log \sum_{z_i \in \{0,1\}} q(z_i) \frac{p(x_i, y_i, z_i|C)}{q(z_i)}, \quad (6)$$

where $q(z_i) = p(z_i|x_i, y_i, C)$. According to the Bayes theorem and Jensen’s inequality, we have

$$\begin{aligned} NLL(D|C) &= - \sum_D \log \sum_{z_i \in \{0,1\}} q(z_i) p(x_i, y_i|C), \\ &\leq - \sum_D \sum_{z_i \in \{0,1\}} q(z_i) \log p(x_i, y_i|C) \\ &= - \sum_D \sum_{z_i \in \{0,1\}} q(z_i) \log p(y_i|C, x_i) + \text{const}, \end{aligned} \quad (7)$$

where $-\sum_D \sum_{z_i \in \{0,1\}} q(z_i) \log p(y_i|C, x_i)$ is the upper bound of $NLL(D|C)$. Typically, we can adopt the EM algorithm to find the prototypes C that minimize the upper bound by iterating:

E-step: Compute a new estimate of $q(z_i)$ (i.e., clean or noisy) according to prototypes C^{old} from the last iteration:

$$q(z_i) = p(z_i|x_i, y_i, C^{old}). \quad (8)$$

M-step: Find the prototypes C that minimizes the bound:

$$C^{new} = \arg \min_C - \sum_D \sum_{z_i \in \{0,1\}} q(z_i) \log p(y_i|C, x_i) \quad (9)$$

In our method, in order to introduce the “small-loss prior” to provide stronger and more robust supervision signals to the learning of CPC, in the **E-step**, we estimate the distribution of clean or noise of samples, denoted as $q(z'_i)$, via the GMM cleaner instead of $q(z_i)$ in Eq. (8). And consequently, we replace the $q(z_i)$ in Eq. (9) to $q(z'_i)$ and find the prototype C minimize the bound. Next, we provide the justification that the EM algorithm still work by proving that $q(z'_i)$ can be considered as an approximation to $q(z_i)$ in our framework.

In our method, $q(z'_i) = p(z'_i|l(p(\tilde{y}_i|x_i), y_i))$, where $\tilde{y}_i \sim p(\tilde{y}_i|x_i, \theta)$, which is the label predicted by the DNN parameterized by θ . As introduced in section 4.1, in the first stage of each epoch, the CPC’s estimation results $z_i \sim q(z_i)$ are utilized to divide training samples into a clean labeled set $\mathcal{X} = \{(x_i, y_i)|z_i = 0\}$ and a noisy unlabeled set $\mathcal{U} = \{(x_i, y_i)|z_i = 1\}$. Then the parameters of DNNs, which we denote as θ , are optimized using Eq. (11) in the second stage. There exists an optimal θ^* with respected to z_i , with which the softmax output $p(\tilde{y}_i|x_i)$ of DNNs satisfies:

$$\ell(p(\tilde{y}_i|x_i), y_i) = 0, \text{ if } z_i = 0, \text{ otherwise } 1, \quad (10)$$

where $\ell(p(\tilde{y}_i|x_i), y_i)$ is the cross-entropy loss between the network prediction and the annotated label. With these loss values, the subsequent GMM cleaner can easily distinguish samples of \mathcal{X} from samples of \mathcal{U} . In other words, under the optimal θ^* , the estimation of the GMM cleaner would be consistent with the partition of CPC, i.e., $z'_i = z_i$. In practice, in each epoch, we takes the θ optimized to minimize Eq. (11) as an approximation to the optimal θ^* with respect to z_i , and consequently we can get $q(z'_i)$ as an approximation to $q(z_i)$. Therefore, we can see that with the “small loss prior” introduced into the prototype learning, the EM optimization procedure would still work, which guarantees CPC can learn meaningful prototypes and act as an effective cleaner. In appendix B, we further present empirical results to demonstrate the approximation is hold in practice.

5 EXPERIMENTS

5.1 DATASETS AND IMPLEMENTATION DETAILS

Datasets. We evaluate our method on the following popular LNL benchmarks. For CIFAR-10 and CIFAR-100 (Krizhevsky et al., 2009), we experiment with two types of synthetic noise: symmetric and asymmetric, which are injected into the datasets following the standard setup in (Li et al., 2020a).

Clothing1M (Xiao et al., 2015b) and WebVision1.0 (Li et al., 2017a) are two large-scale real-world label noise benchmarks. Clothing1M contains 1 million images in 14 categories acquired from online shopping websites, which is heavily imbalanced and most of the noise is asymmetric (Yi & Wu, 2019b). WebVision1.0 contains 2.4 million images crawled from the web using the concepts in ImageNet-ILSVRC12 (ILSVRC12). Following convention, we compare with SOTAs on the first 50 classes of WebVision, as well as the performance after transferring to ILSVRC12.

Implementation details. We plug the proposed CPC to the DivideMix (Li et al., 2020a) framework. For Clothing1M and CIFAR-10 with asymmetric noise, we employ a single class-agnostic GMM for loss-distribution modeling. For other cases, we find that class-aware GMMs would further improve the performance of CPC. Following DivideMix, we employ ResNet18 (He et al., 2016) for CIFAR-10 and CIFAR-100, and utilize ImageNet pre-trained ResNet-50 for Clothing1M. Since previous works chose different backbones, *e.g.*, Inception-resnet v2 (Szegedy et al., 2017) and ResNet-50, we adopt the weaker one, *i.e.*, ResNet-50 according to (Zheltonozhskii et al., 2021), and train it from scratch for fair comparison. The threshold of CPC τ is set 0.5 by default for all the datasets except for the extremely imbalanced Clothing1M where it is set to 0.3. For CIFAR-10 and CIFAR-100, we train the models for 450 epochs. For the large-scale dataset Clothing1M and WebVision1.0, we train the model for 80 and 100 epochs, respectively. The warm-up periods of prototypes for all the datasets is set to the first 5% epochs after network warm-up, except in CIFAR-100 with noise ratios larger than 80% when set to 10% of total epochs. For the other settings, we simply follow the standard set-up as in DivideMix. For more implementation details, please refer to the appendix A.2 and codes in supplementary materials.

Table 1: Comparison with SOTAs on Real-world Benchmarks. Following GJS(Engleson & Azizpour, 2021), we run our method three times with different random seeds and report the mean and standard deviation. \diamond indicates methods utilize ResNet50 for WebVision, while others utilize Inception-resnet v2.

| | WebVision | | WebVision \rightarrow ILSVRC12 | | Clothing1M |
|----------------------|------------------|-------------------------|----------------------------------|-------------------------|-------------------------|
| | top1 | top5 | top1 | top5 | |
| ELR+ | 77.78 | 91.64 | 70.29 | 89.76 | 74.8 |
| DivideMix | 77.32 | 91.64 | 75.2 | 90.84 | 74.76 |
| DivideMix \diamond | 76.3 \pm 0.36 | 90.65 \pm 0.16 | 74.42 \pm 0.29 | 91.21 \pm 0.12 | 74.76 |
| LongReMix | 78.92 | 92.32 | - | - | 74.38 |
| NGC | 79.16 | 91.84 | 74.44 | 91.04 | - |
| AugDMix | - | - | - | - | 75.11 |
| NCR \diamond | 80.5 | - | - | - | 74.6 |
| GJS \diamond | 79.28 \pm 0.24 | 91.22 \pm 0.3 | 75.5 \pm 0.17 | 91.27 \pm 0.26 | - |
| Baseline \diamond | 76.3 \pm 0.36 | 90.65 \pm 0.16 | 74.42 \pm 0.29 | 91.21 \pm 0.12 | 74.73 \pm 0.02 |
| Ours \diamond | 79.63 \pm 0.08 | 93.46 \pm 0.10 | 75.75 \pm 0.14 | 93.49 \pm 0.25 | 75.40 \pm 0.10 |

5.2 COMPARISON WITH STATE-OF-THE-ART METHODS

Real-world noise benchmarks. We evaluate our method on real-world large scale data sets, and compare our method with latest SOTA label noise learning methods, including DivideMix(Li et al., 2020a), LongReMix(Cordeiro et al., 2022), NGC(Wu et al., 2021), GJS(Engleson & Azizpour, 2021), ELR+(Liu et al., 2020), AugDMix(Nishi et al., 2021) and NCR(Huang et al., 2021). For WebVision, we measure the top1 and top5 accuracy on WebVision validation set and ImageNet ILSVRC12 validation set. We take ResNet50-based DivideMix (Zheltonozhskii et al., 2021) as baseline. As shown in Table 1, our CPC improves top1 and top5 accuracy over baseline model on WebVision by 3.33% and 2.81%, respectively. Our method achieves competitive performance on WebVision, and shows stronger transferable capability, outperforming other competitors on the ILSVRC12 validation set significantly. For Clothing1M, we apply the strong augmentation strategy (Nishi et al., 2021) to DivideMix as our baseline, and rerun the method three times. Our method achieves 75.4% accuracy on this challenging benchmark, outperforming all the other SOTAs. We also notice that though NCR achieves SOTA result on WebVision, it shows moderate performance compared to ELR+, DivideMix and AugDMix on Clothing1M containing asymmetric noise with imbalanced data distribution. It reveals that our method could be more robust across different label noise scenarios.

Table 2: Comparison with SOTAs on Synthetic Noise Benchmarks: CIFAR-10 and CIFAR-100. Following previous work (Engleson & Azizpour, 2021; Wu et al., 2021; Zheltonozhskii et al., 2021), we run our method three times with different random seeds and report the mean and standard deviation. \dagger indicates the baseline of our method. * indicates semi-supervised learning based label noise learning methods.

| | CIFAR-10 / CIFAR-100 (Sym) | | | | CIFAR-10 (Asym) |
|--------------------|---|--|--|--|----------------------------|
| | 20% | 50% | 80% | 90% | 40% |
| ELR+ | 94.9 / 76.3 | 93.9 / 72.0 | 90.9 / 57.2 | 74.5 / 30.9 | 88.9 |
| NCR | 95.2 / 76.6 | 94.3 / 72.5 | 91.6 / 58.0 | 75.1 / 30.8 | 90.7 |
| ProtoMix | 96.4 / 80.3 | 95.3 / 76.0 | 93.3 / 61.1 | 77.4 / 33.1 | 92.6 |
| DivideMix* | 96.1 / 77.3 | 94.6 / 74.5 | 93.2 / 60.2 | 76.0 / 31.5 | 93.4 |
| LongReMix* | 96.2 / 77.8 | 95.0 / 75.6 | 93.9 / 62.9 | 82.0 / 33.8 | 94.7 |
| AugDMix* \dagger | 96.3 / 79.5 | 95.6 / 77.2 | 93.6 / 66.4 | 91.9 / 41.2 | 94.6 |
| <hr/> | | | | | |
| NGC | 95.88 / 79.31 ± 0.13 / ± 0.35 | 94.54 / 75.91 ± 0.35 / ± 0.39 | 91.59 / 62.7 ± 0.31 / ± 0.37 | 80.46 / 29.76 ± 1.97 / ± 0.85 | 90.55 ± 0.29 |
| GJS | 95.33 / 75.71 ± 0.18 / ± 0.25 | - | 79.11 / 44.49 ± 0.31 / ± 0.53 | - | 89.65 ± 0.37 |
| Ours* | 96.50 / 80.22 ± 0.10 / ± 0.21 | 95.64 / 79.31 ± 0.01 / ± 0.13 | 94.78 / 69.56 ± 0.01 / ± 0.34 | 92.55 / 54.60 ± 0.58 / ± 0.24 | 94.73 ± 0.04 |

Synthetic noise benchmarks. For better understanding of our work, we also evaluate the performance of CPC on CIFAR-10 and CIFAR-100 datasets with symmetric label noise level ranging from 20% to 90% and asymmetric noise of rate 40%. We take AugDMix as the baseline, and compare our method with latest SOTA methods, where DivideMix, LongReMix and Aug-DMix are semi-supervised learning based methods. Following NGC and GJS, we run our method three times with different random seeds and report the mean and standard deviation. For other methods, *e.g.*, ProtoMix (Li et al., 2021), we report the best results reported in their papers. As shown in Table 2, though with a baseline method as strong as AugDMix, our method brings about performance improvement across all noise levels as well as noise types consistently, and establishes new SOTAs on CIFAR-10 and CIFAR-100. Moreover, CPC shows outstanding performance especially in an extremely noisy condition, which improves the test accuracy on CIFAR-100 at a noise rate of 90% by 13%. Additionally, we notice that, under asymmetric noise set-up, semi-supervised learning based methods consistently outperform other methods that achieve impressive results on WebVision benchmark, including NGC, GJS and NCR. The results reveal that semi-supervised learning based method could be more robust to asymmetric noise, while our method achieves SOTA performance among them.

5.3 ANALYSIS

Is CPC a better label noise cleaner? To evaluate the performance of label noise cleaner, we conduct experiments under both symmetric and asymmetric label noise set-up. For symmetric noise, we use CIFAR-100 with 90% noise as benchmark to figure out why CPC brings such a significant performance improvement under this set-up. For asymmetric noise, we employ the most commonly adopted CIFAR-10-asym40% as benchmark. The AUC of clean/noise binary classification results of a cleaner is calculated as the evaluation metric. We take the original class-agnostic GMM cleaner (GMM_{agn}) proposed in DivideMix as baseline, and compare it to our CPC and the aforementioned naive class-aware GMM cleaner (GMM_{awr}). Furthermore, we also implement another version of CPC that trained based on the class-aware GMM cleaner. To distinguish these two CPC, we denote the regular one trained based on conventional class-agnostic GMM cleaner as CPC_{agn} , and the other one as CPC_{awr} . As shown in Figure 3, in both cases, the regular CPC_{agn} outperforms the baseline GMM_{agn} as well as GMM_{awr} , which demonstrates our class prototype-based method is the better label noise cleaner. As for the comparison between GMM_{agn} and GMM_{awr} , we find that in the situation of high symmetric noise, though GMM_{agn} shows better performance in the early stage of training, GMM_{awr} outperforms it in the second half stage of training. In the case of asymmetric noise, GMM_{awr} , which tend to classify hard clean samples in clean categories as noise wrongly, consistently underperforms GMM_{agn} across the whole training period. The results further prove that our class prototype-based method is the better choice for applying class-aware modulation to




Figure 3: AUC of different label noise cleaners with respect to the training period. Left is the results on CIFAR-100 under high symmetric noise ratio 0.9. Right is the results on CIFAR-10 under medium asymmetric noise ratio 0.4.

Table 3: The effect of different label noise cleaner.

| Model | CIFAR-100 sym-90% | WebVision top1 | CIFAR-10 asym-40% | Clothing1M |
|----------------|----------------------|-------------------|----------------------|-------------|
| w/ GMM_{agn} | 41.2 | 76.32 | 94.6 | 74.73 |
| w/ GMM_{awr} | 45.6 | 78.66 | 94.18 | 74.18 |
| w/ CPC_{agn} | 48.88 | 79.4 | 94.73 | 75.4 |
| w/ CPC_{awr} | 54.6 | 79.63 | 94.29 | 74.36 |

label noise cleaning, which is more robust across different noise types. Moreover, we find that in the case of asymmetric noise, CPC_{agn} achieves higher AUC compared to GMM_{agn} , which shows our method can partially make up for the shortcomings of GMM_{agn} . In the case of symmetric noise, we find that GMM_{agn} can further improve the performance of CPC, where CPC_{awr} achieves the best performance among the four cleaners.

How do different label noise cleaners affect label noise learning? We plug different cleaners to DivideMix framework, and keep all the other training settings the same as described in the implementation details. As shown in Table 3, the final performance of the model is consistent with the performance of the cleaner used. On CIFAR-100 with 90% symmetric noise, performance improvement brought about by CPC_{agn} are 7.68%, while model with CPC_{awr} outperforms the baseline method by 13.4%. We also report the comparison results on large-scale WebVision dataset, where the performance of different models show the same trend of change as in CIFAR-100-sym90%. As for the asymmetric noise situation, *i.e.*, CIFAR-10-asym40% and Clothing1M, model with CPC_{agn} , which has superior label noise partitioning capability as shown in Fig.3, achieves best performance while CPC_{awr} beat GMM_{awr} in both cases. The results demonstrate that CPC is helpful to train a better model in label noise learning.

How about the quality of prototypes learnt in CPC? To evaluate the quality of prototypes learnt in CPC, we utilize the prototypes to classify test samples. We calculate the distance between each class prototype and the embedding of a given sample, and take the class whose prototype has the shortest distance to the sample as the result. We compare our method with another prototype-based label noise learning method MoPro (Li et al., 2020b) on the first 50 classes of WebVision. We implement the experiment with the official code released by the MoPro team. The results show that our prototype achieves a top1 accuracy of 78.44%, while MoPro’s accuracy is 72.23%. The result demonstrates that our method is able to learn better prototypes.

6 CONCLUSION

In this paper, we reveal the long-ignored problem of *class-agnostic* loss distribution modeling that widely existed in label noise learning, and propose a simple yet effective solution, named Class Prototype-based label noise Cleaner (CPC). CPC takes advantage of loss distribution modeling and intra-class consistency regularization in feature space simultaneously, which can better distinguish clean and noise labels. We justify the effectiveness of our method by explaining it from the EM algorithm perspective theoretically and providing extensive empirical proves. The experimental results show that our method achieves competitive performance compared to current SOTAs.

REFERENCES

- Devansh Arpit, Stanisław Jastrzebski, Nicolas Ballas, David Krueger, Emmanuel Bengio, Maxinder S Kanwal, Tegan Maharaj, Asja Fischer, Aaron Courville, Yoshua Bengio, et al. A closer look at memorization in deep networks. In *International conference on machine learning*, pp. 233–242. PMLR, 2017.
- David Berthelot, Nicholas Carlini, Ian Goodfellow, Nicolas Papernot, Avital Oliver, and Colin A Raffel. Mixmatch: A holistic approach to semi-supervised learning. *Advances in neural information processing systems*, 32, 2019.
- Filipe R. Cordeiro, Vasileios Belagiannis, Ian D. Reid, and G. Carneiro. Propmix: Hard sample filtering and proportional mixup for learning with noisy labels. *BMVC*, 2021.
- Filipe R Cordeiro, Ragav Sachdeva, Vasileios Belagiannis, Ian Reid, and Gustavo Carneiro. Longremix: Robust learning with high confidence samples in a noisy label environment. *Pattern Recognition*, pp. 109013, 2022.
- Yifan Ding, Liqiang Wang, Deliang Fan, and Boqing Gong. A semi-supervised two-stage approach to learning from noisy labels. *2018 IEEE Winter Conference on Applications of Computer Vision (WACV)*, pp. 1215–1224, 2018.
- Erik Englesson and Hossein Azizpour. Generalized jensen-shannon divergence loss for learning with noisy labels. *Advances in Neural Information Processing Systems*, 34:30284–30297, 2021.
- Bo Han, Quanming Yao, Xingrui Yu, Gang Niu, Miao Xu, Weihua Hu, Ivor Wai-Hung Tsang, and Masashi Sugiyama. Co-teaching: Robust training of deep neural networks with extremely noisy labels. In *NeurIPS*, 2018.
- Bo Han, Quanming Yao, Tongliang Liu, Gang Niu, Ivor W Tsang, James T Kwok, and Masashi Sugiyama. A survey of label-noise representation learning: Past, present and future. *arXiv preprint arXiv:2011.04406*, 2020.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Identity mappings in deep residual networks. In *European conference on computer vision*, pp. 630–645. Springer, 2016.
- Zhenyu Huang, Guocheng Niu, Xiao Liu, Wenbiao Ding, Xinyan Xiao, Hua Wu, and Xi Peng. Learning with noisy correspondence for cross-modal matching. *Advances in Neural Information Processing Systems*, 34:29406–29419, 2021.
- Lu Jiang, Zhengyuan Zhou, Thomas Leung, Li-Jia Li, and Li Fei-Fei. Mentornet: Learning data-driven curriculum for very deep neural networks on corrupted labels. In *ICML*, 2018.
- Lu Jiang, Di Huang, Mason Liu, and Weilong Yang. Beyond synthetic noise: Deep learning on controlled noisy labels. In *ICML*, 2020.
- Alex Krizhevsky, Geoffrey Hinton, et al. Learning multiple layers of features from tiny images. 2009.
- Kuang-Huei Lee, Xiaodong He, Lei Zhang, and Linjun Yang. Cleannet: Transfer learning for scalable image classifier training with label noise. *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 5447–5456, 2018.
- Junnan Li, Richard Socher, and Steven C. H. Hoi. Dividemix: Learning with noisy labels as semi-supervised learning. *ICLR*, 2020a.
- Junnan Li, Caiming Xiong, and Steven CH Hoi. Mopro: Webly supervised learning with momentum prototypes. *arXiv preprint arXiv:2009.07995*, 2020b.
- Junnan Li, Pan Zhou, Caiming Xiong, and Steven CH Hoi. Prototypical contrastive learning of unsupervised representations. *arXiv preprint arXiv:2005.04966*, 2020c.
- Junnan Li, Caiming Xiong, and Steven CH Hoi. Learning from noisy data with robust representation learning. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 9485–9494, 2021.

- Wen Li, Limin Wang, Wei Li, Eirikur Agustsson, and Luc Van Gool. Webvision database: Visual learning and understanding from web data. *arXiv preprint arXiv:1708.02862*, 2017a.
- Yuncheng Li, Jianchao Yang, Yale Song, Liangliang Cao, Jiebo Luo, and Li-Jia Li. Learning from noisy labels with distillation. *2017 IEEE International Conference on Computer Vision (ICCV)*, pp. 1928–1936, 2017b.
- Sheng Liu, Jonathan Niles-Weed, Narges Razavian, and Carlos Fernandez-Granda. Early-learning regularization prevents memorization of noisy labels. *Advances in neural information processing systems*, 33:20331–20342, 2020.
- Frank J Massey Jr. The kolmogorov-smirnov test for goodness of fit. *Journal of the American statistical Association*, 46(253):68–78, 1951.
- Kento Nishi, Yi Ding, Alex Rich, and Tobias Hollerer. Augmentation strategies for learning with noisy labels. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 8022–8031, 2021.
- Scott Reed, Honglak Lee, Dragomir Anguelov, Christian Szegedy, Dumitru Erhan, and Andrew Rabinovich. Training deep neural networks on noisy labels with bootstrapping. *arXiv preprint arXiv:1412.6596*, 2014.
- Christian Szegedy, Sergey Ioffe, Vincent Vanhoucke, and Alexander A Alemi. Inception-v4, inception-resnet and the impact of residual connections on learning. In *Thirty-first AAAI conference on artificial intelligence*, 2017.
- Daiki Tanaka, Daiki Ikami, T. Yamasaki, and Kiyoharu Aizawa. Joint optimization framework for learning with noisy labels. *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 5552–5560, 2018.
- Arash Vahdat. Toward robustness against label noise in training deep discriminative neural networks. In *NIPS*, 2017.
- Haobo Wang, Ruixuan Xiao, Yixuan Li, Lei Feng, Gang Niu, Gang Chen, and Junbo Zhao. Pico: Contrastive label disambiguation for partial label learning. *arXiv preprint arXiv:2201.08984*, 2022.
- Zhi-Fan Wu, Tong Wei, Jianwen Jiang, Chaojie Mao, Mingqian Tang, and Yu-Feng Li. Ngc: a unified framework for learning with open-world noisy data. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 62–71, 2021.
- Tong Xiao, Tian Xia, Yi Yang, Chang Huang, and Xiaogang Wang. Learning from massive noisy labeled data for image classification. *2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 2691–2699, 2015a.
- Tong Xiao, Tian Xia, Yi Yang, Chang Huang, and Xiaogang Wang. Learning from massive noisy labeled data for image classification. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 2691–2699, 2015b.
- Kun Yi and Jianxin Wu. Probabilistic end-to-end noise correction for learning with noisy labels. *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 7010–7018, 2019a.
- Kun Yi and Jianxin Wu. Probabilistic end-to-end noise correction for learning with noisy labels. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 7017–7025, 2019b.
- Xingrui Yu, Bo Han, Jiangchao Yao, Gang Niu, Ivor Wai-Hung Tsang, and Masashi Sugiyama. How does disagreement help generalization against label corruption? In *ICML*, 2019.
- Hongyi Zhang, Moustapha Cisse, Yann N Dauphin, and David Lopez-Paz. mixup: Beyond empirical risk minimization. *arXiv preprint arXiv:1710.09412*, 2017.
- Evgenii Zheltonozhskii, Chaim Baskin, Avi Mendelson, Alex M. Bronstein, and Or Litany. Contrast to divide: Self-supervised pre-training for learning with noisy labels. *CoRR*, abs/2103.13646, 2021. URL <https://arxiv.org/abs/2103.13646>.

A APPENDIX

A.1 EMPIRICAL VICINAL RISK

We introduce the Empirical Vicinal Risk following [Cordeiro et al. \(2022\)](#). In the semi-supervised learning based label noise learning framework, with the clean labeled set \mathcal{X} and noisy unlabeled set \mathcal{U} from a cleaner, the DNNs are trained to minimise the empirical vicinal risk (EVR) ([Zhang et al., 2017](#)):

$$\ell_{EVR} = \frac{1}{|\mathcal{X}'|} \sum_{(x'_i, y'_i) \in \mathcal{X}'} \ell_{\mathcal{X}'}(p(\tilde{y}'_i|x'_i), y'_i) + \frac{\lambda^{(\mathcal{U}')}}{|\mathcal{U}'|} \sum_{(x'_i, y'_i) \in \mathcal{U}'} \ell_{\mathcal{U}'}(p(\tilde{y}'_i|x'_i), y'_i), \quad (11)$$

where $\ell_{\mathcal{X}'}$ and $\ell_{\mathcal{U}'}$ denote the losses for set \mathcal{X}' and \mathcal{U}' , which are weighted by $\lambda^{(\mathcal{U}'})$. \mathcal{X}' and \mathcal{U}' indicate MixMatch ([Berthelot et al., 2019](#)) augmented clean and noise set:

$$\begin{aligned} \mathcal{X}' &= (x'_i, y'_i) : (x'_i, y'_i) \sim f(x'_i, y'_i|x_i, y_i), (x_i, y_i) \in \mathcal{X}, \\ \mathcal{U}' &= (x'_i, y'_i) : (x'_i, y'_i) \sim f(x'_i, y'_i|x_i, y_i), (x_i, y_i) \in \mathcal{U}, \end{aligned} \quad (12)$$

with

$$f(x'_i, y'_i|x_i, y_i) = \frac{1}{|\mathcal{X} \cup \mathcal{U}|} \sum_{(x_i, y_i) \in \mathcal{X} \cup \mathcal{U}} \mathbb{E}_\lambda[\delta(x'_i = x_i + (1 - \lambda)x_j, y'_i = y_i + (1 - \lambda)y_j)], \quad (13)$$

where δ is a Dirac mass centered at (x', y') , $\lambda \sim \text{Beta}(a, a)$, and $a \in (0, \inf)$.

A.2 TRAINING CONFIGURES

In our method, we follow most of training set-up of DivideMix([Li et al., 2020a](#)). We present the detailed training configures as follows:

- **CIFAR-10 and CIFAR-100.** For all the experiments on CIFAR, we train our DNN model as well as class prototypes in CPC via SGD with a momentum of 0.9, a weight decay of 0.0005, and a batch size of 128. The network is trained for 450 epochs. We set the initial learning rate as 0.02, and reduce it by a factor of 10 after 225 epochs. The warm up period for the DNN is 10 epochs. The weight $\lambda^{(\mathcal{U}')}$ is set to $\{0, 25, 50, 150\}$ as in DivideMix.
- **Clthing1M.** We train our DNN model as well as class prototypes in CPC via SGD with a momentum of 0.9, a weight decay of 0.001, and a batch size of 32. The model is trained for 80 epochs. The warm up period for the DNN is 1 epoch. The initial learning rate is set as 0.002 and reduced by a factor of 10 after 40 epochs. For each epoch, we sample 1000 mini-batches from the training data. The weight $\lambda^{(\mathcal{U}')}$ is set to 0.
- **WebVision.** We train our DNN model as well as class prototypes in CPC via SGD with a momentum of 0.9, a weight decay of 0.001, and a batch size of 32. The model is trained for 100 epochs. The warm up period for the DNN is 1 epoch. The initial learning rate is set as 0.01 and reduced by a factor of 10 after 50 epochs. For each epoch, we sample 1000 mini-batches from the training data. The weight $\lambda^{(\mathcal{U}')}$ is set to 0.

B IS $q(z'_i)$ A PROPER APPROXIMATION TO $q(z_i)$ IN PRACTICAL?

In Section 4.3, we replace the estimation of CPC $q(z_i)$ in Eq. (9) with the estimation of GMM cleaner $q(z'_i)$ and justify $q(z'_i)$ can be considered as an approximate to $q(z_i)$. To investigate if the approximation holds in practical, we calculate the K-L Divergence as well as classification consistency between $q(z'_i)$ and $q(z_i)$. As shown in Figure 4, as the training going on, the KLD between $q(z'_i)$ and $q(z_i)$ is converged and the classification consistency increases.




Figure 4: The left figure shows the KLD between $q(z'_i)$ and $q(z_i)$. The right figure presents the consistency rate between z'_i and z_i . Results are collected from CIFAR-10-aysm40%.